

# CultiVar-17: One-Shot Digit Recognition with a Culturally-Motivated Style Taxonomy

Attila Torda  
DigitRecognizer Project  
[affiliation placeholder]

**Abstract**—We introduce CultiVar-17, a compact dataset of 17 hand-drawn digit templates that captures stylistic variation motivated by regional and educational handwriting conventions—crossed vs. uncrossed zero, European vs. American seven, serif vs. plain one, and more. Using only these 17 templates as supervision, we train a CNN and a prototype-embedding model via synthetic augmentation and evaluate transfer to the MNIST test set. Our best configuration—a prototype-embedding model with full augmentation—achieves  $77.46 \pm 2.39\%$  MNIST accuracy from 17 examples, reaching 78% of supervised CNN performance. An ablation study shows that elastic distortion and stroke-width perturbation contribute independently and additively. All code and data are publicly available for reproducibility.

**Index Terms**—one-shot learning, digit recognition, handwriting style, data augmentation, prototype networks, MNIST

## I. INTRODUCTION

Learning from very few examples is a longstanding challenge in machine learning [1]–[3]. Handwritten digit recognition has served as a foundational benchmark since LeCun et al. [4] established MNIST, yet standard pipelines assume thousands of labelled samples per digit class.

This work explores a strict one-shot regime: we build a complete recognition pipeline from exactly **one hand-drawn template per class**, structured around a culturally-aware 17-class taxonomy. Our contributions are threefold:

- 1) **CultiVar-17**—a 17-template dataset encoding documented regional writing variants, distributed as a reproducible `.npy` artifact.
- 2) **A one-shot augmentation pipeline**—elastic distortion, stroke-width perturbation, and affine noise applied to templates, generating diverse training batches from a single source.
- 3) **An empirical study**—comparing nearest-template, CNN, and prototype-embedding approaches with ablation over augmentation components, achieving 77.46% MNIST accuracy from 17 examples.

## II. RELATED WORK

### A. One-Shot and Few-Shot Learning

Koch et al. [1] proposed Siamese Networks for one-shot image recognition, learning a similarity metric between pairs of examples. Vinyals et al. [2] introduced Matching Networks, which use attention-weighted nearest-neighbour classification in an embedding space conditioned on the support set. Snell et al. [3] proposed Prototypical Networks, classifying by

TABLE I  
CULTURAL MOTIVATION FOR CULTIVAR-17 VARIANT SPLITS.

| Digit | Var. A      | Var. B                 | Driver                |
|-------|-------------|------------------------|-----------------------|
| 0     | uncrossed   | crossed                | zero-slash convention |
| 1     | plain vert. | serif / angled         | American vs. European |
| 2     | looped base | angular base           | cursive vs. print     |
| 4     | open top    | closed top (+ cursive) | printing vs. cursive  |
| 7     | uncrossed   | crossed                | American vs. European |
| 9     | looped tail | straight tail          | print vs. cursive     |

computing Euclidean distances to class prototype embeddings—the architecture we adopt here. Finn et al. [5] introduced MAML, a gradient-based meta-learning approach that learns an initialisation enabling fast adaptation with few gradient steps.

### B. Handwriting Datasets and Cultural Variation

Lake et al. [6] introduced Omniglot, a 1,623-character handwriting dataset covering 50 alphabets from diverse writing systems, which became the canonical few-shot classification benchmark. CultiVar-17 is complementary in scope: rather than cross-alphabet generalisation, it targets intra-digit stylistic variation driven by documented regional conventions. LeCun et al. [4] established MNIST as the standard digit recognition benchmark; we use MNIST as the transfer target throughout.

### C. Data Augmentation for Handwriting

Simard et al. [7] demonstrated that elastic distortions—smooth random displacement fields applied to handwritten characters—substantially improve CNN generalisation on MNIST. We apply the same elastic distortion technique as a core augmentation component, combined with stroke-width perturbation. Our ablation (Section VI-B) confirms their independent and additive contributions.

## III. CULTIVAR-17 DATASET

### A. Cultural Motivation

Handwriting style reflects regional education systems and typographic conventions. We document seven split points in standard digit classes, shown in Table I. Digits 3, 5, 6, and 8 have no variant split, yielding  $6 \times 2 + 1 \times 3 + 4 = 17$  classes in total.

### B. Template Extraction

Source: a single composite image generated with an AI image tool and manually verified for legibility and distinctiveness. Procedure: isolate the central strip; split into 17 equal horizontal slots; foreground-tight crop and centre each digit onto a  $28 \times 28$  canvas. Figure 1 shows all 17 templates.

Output: `train_images.npy` ( $17 \times 28 \times 28$  uint8) and `train_labels17.npy` (labels 0–16).

### C. Class-to-Digit Mapping

For MNIST evaluation, 17 style classes are collapsed to 10 canonical digits via a fixed surjective map (`src/variants17/label_schema.py`). Each variant maps to its base digit (e.g., `0_variant_a` and `0_variant_b`  $\rightarrow$  0).

## IV. METHOD

### A. Augmentation Pipeline

Given one template per class, we generate training batches via stochastic augmentation:

- **Elastic distortion** [7]: random displacement field ( $\sigma=4$ ,  $\alpha \in [15, 32]$ ), interpolated with Gaussian smoothing.
- **Stroke-width perturbation**: morphological dilation (`disk_dilate`, radius  $\in \{1, 2\}$ ) or soft blur (`blur_soft`), applied with probability 0.80 per sample.
- **Affine noise**: rotation  $\pm 8^\circ$ , scale  $\in [0.92, 1.08]$  per axis.
- **Additive noise**: Gaussian, std  $\in [0.01, 0.03]$ .

With 256 repeats per class, each epoch presents 4,352 synthetically varied examples. Figure 2 illustrates the six augmentation modes.

### B. Models

**CNN (SimpleCNN)**. Two convolutional blocks (conv  $\rightarrow$  BN  $\rightarrow$  ReLU  $\rightarrow$  pool) followed by a fully-connected head with 17-way softmax.

**Prototype embedding (EmbeddingCNN)**. The same convolutional backbone projects each image into a 64-dimensional embedding. Prototypes are the mean embeddings of the 17 raw templates. At test time a query is classified by nearest-prototype distance [3].

Both models use Adam ( $\text{lr} = 10^{-3}$ , batch 128) for 8 epochs.

## V. EXPERIMENTS

### A. Configurations

Table II lists the five experimental configurations. CNN configurations run for 3 seeds; the prototype model for 5 seeds. Best-epoch MNIST accuracy is reported per seed.

### B. Baseline

**Nearest-template (L2)** [8]: assign each MNIST image the digit label of its closest template under Euclidean distance. No training is required.

TABLE II  
EXPERIMENTAL CONFIGURATIONS.

| Config            | Elastic prob | Stroke prob |
|-------------------|--------------|-------------|
| no_aug_cnn        | 0.00         | 0.00        |
| elastic_only_cnn  | 0.70         | 0.00        |
| stroke_only_cnn   | 0.00         | 0.80        |
| full_aug_cnn      | 0.70         | 0.80        |
| proto (embedding) | 0.70         | 0.80        |

TABLE III  
MNIST TEST ACCURACY (MEAN  $\pm$  STD %).

| Method                     | Acc. (%)                           | Seeds |
|----------------------------|------------------------------------|-------|
| Nearest-template (L2)      | 38.60                              | —     |
| No-aug CNN                 | $51.20 \pm 1.86$                   | 3     |
| Full-aug CNN               | $65.19 \pm 2.14$                   | 3     |
| <b>Proto embedding</b>     | <b><math>77.46 \pm 2.39</math></b> | 5     |
| Supervised CNN (reference) | 99.11                              | —     |

## VI. RESULTS

### A. Main Results

Table III reports MNIST test-set accuracy for all methods. Figure 3 visualises the comparison.

The prototype-embedding model reaches 77.46%—**78% of supervised CNN performance**—from 17 templates. Even the no-augmentation CNN (51.20%) substantially outperforms the nearest-template baseline (38.60%), showing that training on augmented templates provides significant discriminative capacity beyond simple template matching.

### B. Augmentation Ablation

Table IV reports the  $2 \times 2$  ablation of elastic distortion and stroke-width perturbation.

Both components independently improve accuracy by  $\sim 5$  percentage points over no augmentation, and their combination yields a further gain ( $\sim 9$  pp over elastic alone). The higher variance of *elastic\_only* ( $\pm 4.44$ ) versus *stroke\_only* ( $\pm 0.78$ ) suggests elastic distortion is more seed-sensitive, while stroke perturbation provides more consistent improvement.

## VII. DISCUSSION

Three patterns emerge:

**Augmentation is essential.** The gap between nearest-template (38.60%) and full-aug CNN (65.19%) confirms that synthetic diversity substantially improves generalisation beyond direct pixel matching.

**Embedding beats classification.** The prototype model (77.46%) outperforms the classification CNN (65.19%) despite using the same backbone, consistent with findings from Snell et al. [3] that the metric-learning objective is better matched to the few-shot regime.

**The ceiling is style coverage, not capacity.** Even at 77.46%, the system falls 22 pp short of supervised accuracy. This gap reflects the fact that 17 templates cannot cover all handwriting styles in MNIST regardless of augmentation or model complexity.

TABLE IV  
AUGMENTATION ABLATION (CNN, 3 SEEDS, MNIST TEST %).

| Augmentation            | Elastic | Stroke | Acc. (%)     |
|-------------------------|---------|--------|--------------|
| None                    | ✗       | ✗      | 51.20 ± 1.86 |
| Elastic only            | ✓       | ✗      | 56.79 ± 4.44 |
| Stroke only             | ✗       | ✓      | 55.59 ± 0.78 |
| Elastic + stroke (full) | ✓       | ✓      | 65.19 ± 2.14 |

#### A. Augmentation Realism vs. Accuracy

An unexpected finding emerged from a post-hoc augmentation ablation. The stroke-width dilation uses morphological dilation with radius  $\in \{1, 2\}$  and a hard intensity floor of 0.85, producing visually unrealistic images—thick, near-binary strokes resembling ink blobs rather than handwritten digits. Three attempts to replace this with perceptually natural alternatives were tested:

TABLE V  
PROTO ACCURACY UNDER ALTERNATIVE AUGMENTATION VARIANTS.

| Augmentation variant                      | Proto acc. (%)      |
|-------------------------------------------|---------------------|
| Original (radius 1–2, floor 0.85)         | <b>77.46 ± 2.39</b> |
| Softer (radius 1, floor 0.65)             | 75.55 ± 2.39        |
| Softer + edge smoothing                   | 75.18 ± 0.85        |
| Wider Gaussian blur ( $\sigma \leq 2.0$ ) | 75.21 ± 2.26        |

Every visually improved variant costs approximately 2 pp. The explanation is distributional: MNIST contains a sub-population of heavy-handed writers whose thick, dense strokes resemble the blob augmentation. Removing the blobs reduces coverage of this sub-population. This illustrates a general principle: *augmentation realism and benchmark coverage are not the same objective*. The original augmentation is retained because it maximises the reported metric; systems trained on real handwriting samples would not require it.

#### VIII. THREATS TO VALIDITY

- **Single-source template bias:** all 17 templates originate from one AI-generated composite image and reflect a single stylistic hand.
- **Limited augmentation family:** elastic and stroke perturbations may under-represent real handwriting variability (pen angle, slant, baseline drift).
- **Projection artefacts:** the 17→10 class collapse may hide class-specific errors; a full 17-way confusion matrix would give a more complete picture.

#### IX. FUTURE WORK

Adding even 2–3 real handwriting samples per style cluster would substantially broaden template coverage. Stronger augmentation (local nonlinear warps, pen-angle simulation) and meta-learning objectives such as MAML [5] or Matching Networks [2] could further improve generalisation. Extending CultiVar-17 to additional cultural conventions and non-Latin scripts is a natural next step for the dataset contribution. Finally, the current 8-epoch training budget is conservative;




Fig. 1. All 17 CultiVar templates (8× upscaled). Class labels below each digit indicate the style variant.

increasing to 16–24 epochs is expected to recover accuracy gaps introduced by softer augmentation variants without any augmentation changes.

#### X. REPRODUCIBILITY

All experiments are reproducible with a single command:

```
python scripts/run_one_shot_experiment.py \
    --ablation --proto-seeds 5
```

Figures are regenerated with:

```
python scripts/make_figures.py
```

Raw results are stored in `experiments/reports/one_shot_results.json`.

#### ACKNOWLEDGEMENTS

The author thanks the open-source community for the tools used in this work, in particular the NumPy, PyTorch, scikit-image, and Streamlit projects. The MNIST dataset was provided by Yann LeCun and colleagues and is freely available at <http://yann.lecun.com/exdb/mnist/>.

#### REFERENCES

- [1] G. Koch, R. Zemel, and R. Salakhutdinov, “Siamese neural networks for one-shot image recognition,” in *ICML Deep Learning Workshop*, 2015.
- [2] O. Vinyals, C. Blundell, T. Lillicrap, K. Kavukcuoglu, and D. Wierstra, “Matching networks for one shot learning,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2016.
- [3] J. Snell, K. Swersky, and R. Zemel, “Prototypical networks for few-shot learning,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [4] Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner, “Gradient-based learning applied to document recognition,” *Proceedings of the IEEE*, vol. 86, no. 11, pp. 2278–2324, 1998.
- [5] C. Finn, P. Abbeel, and S. Levine, “Model-agnostic meta-learning for fast adaptation of deep networks,” in *International Conference on Machine Learning (ICML)*, 2017.
- [6] B. M. Lake, R. Salakhutdinov, and J. B. Tenenbaum, “Human-level concept learning through probabilistic program induction,” *Science*, vol. 350, no. 6266, pp. 1332–1338, 2015.
- [7] P. Y. Simard, D. Steinkraus, and J. C. Platt, “Best practices for convolutional neural networks applied to visual document analysis,” in *International Conference on Document Analysis and Recognition (ICDAR)*, 2003.
- [8] T. Cover and P. Hart, “Nearest neighbor pattern classification,” *IEEE Transactions on Information Theory*, vol. 13, no. 1, pp. 21–27, 1967.



Fig. 2. Augmentation gallery for template 0 (4\_variant\_a). Columns left to right: original, elastic (mild), elastic (strong), stroke dilate, stroke blur, full augmentation (5 samples each).




Fig. 3. MNIST test accuracy for all one-shot configurations. Error bars show  $\pm 1$  standard deviation across seeds. Dashed line: supervised CNN reference (99.11%).